

Nikos Kanakaris

Curriculum Vitae

Personal details

<i>Legal name</i>	Nikolaos Kanakaris
<i>Email</i>	nikos.kanakaris89@gmail.com
<i>GitHub</i>	nkanak
<i>Link</i>	nkanak.github.io
<i>Publications</i>	Google scholar
<i>Tel. number</i>	+1 (626) 319-0519
<i>Address</i>	Pasadena 91106, CA

Short bio

- Currently: Postdoctoral research scientist working on graph neural networks and large language models for brain image analysis and systems optimization.
- Ph.D. in machine learning, graph mining, natural language processing. BSc. in Computer Science.
- Experience both in academia and industry.
- Experience in research, end-to-end machine learning projects, full-stack software development and deployment.
- 19 full-paper scientific publications (4 journals, 3 book chapters, 10 conferences, 2 pre-prints)

Professional experience

Postdoctoral research scientist

April 2024 - present

University of Southern California, Full-time

- Heterogeneous computing, code and systems optimization using graph neural networks and large language models. In collaboration with the Intel Corporation Research Lab.
- Interpretable machine learning frameworks to synergize brain age estimation and neuroimaging genetics.

Data Scientist

January 2023 - March 2024

efood - Delivery Hero, Full-time

- Designing and implementing a two-tower neural network model for cross-sell product recommendations during the global Hackathon of Delivery Hero.
- Designing and implementing an ML model for price range prediction for restaurants.
- Designing and implementing a product taxonomy and tagging model using multi-label and multi-task learning techniques. 'item Designing and implementing an in-house product annotation tool.

Technologies: Big Data, Structured and Unstructured data, PyTorch, Hugging Face, BERT, LightGBM, sklearn, pandas, MLflow, Apache Airflow, BigQuery, Python, Ray, Multi-label classification, Seq2Seq models, Ranking, A/B testing, Streamlit, Optuna, Word embeddings, fasttext, SQL

Research Scientist

December 2017 - July 2023

University of Patras, Full-time

- Designing and implementing a novel machine learning pipeline for detecting fake news campaigns using graph convolutional networks for the [Astroturfing.gr](#) research project. In collaboration with the Harokopio University of Athens ([HUA](#))
- Designing and implementing the data analysis pipeline, node classification, node clustering and node representation learning models for the [inPOINT](#) research project
- Designing and implementing the collaboration and discourse analysis environment for [OpenBioC](#) research project
- Working on an information system that monitors smart devices on cruise ships to provide maximum energy savings. [ECLiPse](#) research project
- Implementing the personalization and user profiling algorithms for [MyCorridor](#) European research project

Technologies: Data engineering and data science, crawling twitter data, PyTorch, PyTorch Geometric, text classification, graph classification, representation learning, word embeddings, deep learning, data science, TensorFlow, Node.js, Natural Language Processing (NLP), cluster analysis, discourse analysis, Python, C++, PostgreSQL, MongoDB, Graph Mining, Docker, FastAPI, MongoDB, Neo4j

Software Engineer

February 2016 - August 2017

Programize LLC, Full-time

- Converting an already existing financial model from Microsoft Excel to Python
- Contributing to the automation of the process of converting Microsoft Excel workbooks to Python
- Implementing the back-end system for [Litta](#) mobile application
- Implementing the main mobile hybrid and back-end application for [Global Showcases](#)
- Adding responsive functionality to the official website for [VForVacation](#)
- Contributing to the front-end and back-end internal financial application for [SunPower](#)
- Implementing an Internet Explorer extension for [OnionID](#)

Technologies: Docker, Python, JavaScript, C#, C++, AngularJS, Flask, Node.js, PostgreSQL, Apache Cordova, MySQL, BHO IE, scipy, py2exe

Software Engineer

November 2015 - May 2016

Software Competitiveness International S.A. (SoftCom International), Full-time

Working for a German Automotive Industry under an NDA, area: Navigation

- Implementing the middleware component that facilitates the front-end and back-end software layers

Technologies: C++, Python, Lua, XML

Software and Research Engineer

July 2015 - October 2015

Software Competitiveness International S.A. (SoftCom International), Internship

Working for a Space Industry under an NDA, area: preventive performance monitoring

- [Researching](#) for an alternative algorithm for hierarchical clustering

Technologies: C++, Java, machine learning techniques, data clustering

Education

Ph.D. in Graph Mining and NLP

February 2018 - July 2023

University of Patras

Laboratory: Industrial Management and Information Systems

Specialization: Graph Mining, Natural Language Processing, Text Representations

Dissertation title: *Graph-based knowledge representation and extraction from unstructured textual data using machine learning algorithms* [[Dissertation](#)] [[Presentation](#)]

Description: Exploiting techniques and methods from the fields of graph theory, graph mining, natural language processing, knowledge management and machine learning in order to improve graph-based text representation and graph-based natural language processing. Advancing common natural language processing tasks, such as text classification, feature extraction, feature selection and representation learning, using graph-based text representations, word embeddings, and graph neural networks.

Dissertation Technologies: PyTorch, PyTorch Geometric, Python, Neo4j, TensorFlow, Keras, scikit-learn, pandas

Keywords: Graph mining, representation learning, word embeddings, text representations, text classification, graph neural networks, evaluation and deployment of ML models, data analysis

Other activities: **Activity 1.** Worked on one EU-funded and four national research projects as a researcher where I designed, developed and deployed machine learning models and knowledge management systems. **Activity 2.** Collaborating with industrial companies under an NDA to analyze real-world data, perform market basket analysis and extract insights. **Activity 3.** Teaching assistant. **Activity 4.** Writing proposals and project deliverables. **Activity 5.** Supervision of eight undergraduate and graduate students.

Supervisor: [Nikos Karacapilidis](#)

BSc. Informatics & Telematics

September 2011 - June 2016

Harokopio University of Athens (8.5/10, valedictorian)

Specialization: Computer Software & Systems

Thesis: *Parallelization of “burn scar mapping” algorithms - Παράλληλοποίηση αλγορίθμων χαρτογράφησης καμένων εκτάσεων σε δορυφορικά δεδομένα* [[Thesis](#)] [[Presentation](#)]

Thesis Technologies: Digital Image Processing, Parallel Programming, Python, MPI

Supervisor: [Dimitrios Michail](#)

Skills

Languages: Greek (native), English (fluent)

Programming Languages: C/C++, Python/Cython, JavaScript/Node.js/TypeScript, Julia, Java, SQL, \LaTeX

Databases: PostgreSQL, MySQL, SQLite, Neo4j, MongoDB, BigQuery

Computer Science: Machine Learning, Natural Language Processing, Graph Mining, Graph Neural Networks, Graph Theory, Transformers, LLMs, Recommender Systems, Parallel Programming, Embedded Systems, Full-stack software development

Technologies/Tools/Platforms: PyTorch, PyTorch Geometric, TensorFlow, Keras, scikit-learn, Hugging Face, BERT, LightGBM, Pandas, Docker, Git, NGINX, MLFlow, Optuna, NetworkX, Streamlit, Django, Flask, FastAPI, Express, Angular

Other: Autodesk Fusion 360, 3D printing, Electronics

Interests

graph mining, graph-based text representations, graph neural networks, natural language processing, machine learning, business intelligence, graph theory, algorithms, parallel programming, embedded systems, IoT

Other Activities

Reviewer: [ECML-PKDD](#), [ECAI](#), [AIRE](#), [PMJ](#), [PLOS ONE](#)

Program Committee Member: [ECML-PKDD](#), [ECAI](#), [AIAI](#)

Manuscript/Book Reviewer @ Manning Publications:

- [Go Web Programming, Second Edition](#) by Sau Sheong Chang
- [Julia for Data Analysis](#) by Bogumił Kamiński
- [Julia as a Second Language](#) by Erik Engheim
- [Deep Learning with Python, Second Edition](#) by François Chollet
- [Real-World Natural Language Processing](#) by Masato Hagiwara
- [TensorFlow 2.0 in Action](#) by Thushan Ganegedara
- [Transfer Learning for Natural Language Processing](#) by Paul Azunre
- [Machine Learning for Business](#) by Doug Hudgeon and Richard Nichol
- [Succeeding with AI](#) by Veljko Krunic

Teaching assistant: Semester course ‘Computer Programming’ at the Department of Mechanical Engineering and Aeronautics, University of Patras, Greece (academic years: 2018, 2019, 2020)

Publications

1. S. Duan, H. Ping, N. Kanakaris, X. Xiao, P. Zhang, P. Kyriakis, N. K. Ahmed, G. Ma, M. Capota, T. L., S. Nazarian, Willke, P. Bogdan: A structure-aware framework for learning device placements on computation graphs. arXiv, 2024. ([link](#))
2. S. Duan, N. Kanakaris, X. Xiao, H. Ping, C. Zhou, N. K. Ahmed, G. Ma, M. Capota, T. L. Willke, S. Nazarian, P. Bogdan: Leveraging Reinforcement Learning and Large Language Models for Code Optimization. arXiv, 2023. ([link](#))
3. E. Adamides, N. Giarelis, N. Kanakaris, N. Karacapilidis, K. Konstantinopoulos and I. Siachos: Leveraging Open Innovation Practices Through a Novel ICT Platform. International KES Conference on Human Centred Intelligent Systems, 2023. ([link](#))

4. N. Kanakaris, I. Siachos and N. Karacapilidis: Is it a bug or a feature? Identifying software bugs using graph attention networks. International Conference on Tools with Artificial Intelligence (ICTAI), 2022. ([link](#))
5. N. Kanakaris and N. Karacapilidis: Predicting prices of Airbnb listings via Graph Neural Networks and Document Embeddings: The case of the island of Santorini. International Conference on ENTERprise Information Systems (CENTERIS), 2022. ([link](#))
6. N. Kanakaris, D. Michail and I. Varlamis: A comparative survey of graph databases and software for social network analytics: The link prediction perspective. Book chapter for Graph Databases and their use in social media and smart cities, Science Publishers and CRC Press, Taylor & Francis Group, 2022. ([link](#))
7. D. Michail, N. Kanakaris and I. Varlamis: Detection of fake news campaigns using graph convolutional networks. International Journal of Information Management Data Insights. 2, 100104, 2022. ([link](#))
8. N. Kanakaris, N. Giarelis, I. Siachos and N. Karacapilidis: Making personnel selection smarter through word embeddings: A graph-based approach. Machine Learning with Applications, 100214, 2021, doi: 10.1016/j.mlwa.2021.100214. ([link](#))
9. A. Kanterakis, N. Kanakaris, M. Koutoulakis, K. Pitianou, N. Karacapilidis, L. Koumakis and G. Potamias: Converting Biomedical Text Annotated Resources into FAIR Research Objects with an Open Science Platform. Applied Sciences, Vol. 11, No 20, 9648, 2021, doi: 10.3390/app11209648. ([link](#))
10. N. Kanakaris, N. Giarelis, I. Siachos and N. Karacapilidis: Shall I Work with Them? A Knowledge Graph-Based Approach for Predicting Future Research Collaborations. Entropy, Vol. 23, No 6, 664, 2021, doi: 10.3390/e23060664. ([link](#))
11. N. Giarelis, N. Kanakaris and N. Karacapilidis: Medical Knowledge Graphs in the Discovery of Future Research Collaborations. In: Chee-Peng Lim, Ashlesha Vaidya, Kiran Jain, Virag U Mahorkar, Lakhmi C. Jain (eds.), Handbook of Artificial Intelligence in Healthcare. Springer, 2021 ([link](#)).
12. N. Giarelis, N. Kanakaris and N. Karacapilidis: A comparative assessment of state-of-the-art methods for multilingual unsupervised keyphrase extraction. In: I. Maglogiannis, J. Macintyre and L. Iliadis (eds.), Proceedings of the 17th International Conference on Artificial Intelligence Applications and Innovations (AIAI 2021), Crete, Greece, June 25-27, 2021, IFIP Advances in Information and Communication Technology, Vol. 627, pp. 635-645, doi: 10.1007/978-3-030-79150-6_50.
13. N. Giarelis, N. Kanakaris and N. Karacapilidis: On the utilization of structural and textual information of a scientific knowledge graph to discover future research collaborations: a link prediction perspective. In: A. Appice, G. Tsoumakas, Y. Manolopoulos and S. Matwin (eds.), Proceedings of the 23rd International Conference on Discovery Science (DS 2020), Online Conference, October 19-21, 2020, Springer, Cham, Switzerland, Lecture Notes in Artificial Intelligence, Vol. 12323, pp. 437-450, doi: 10.1007/978-3-030-61527-7_29. ([link](#))
14. N. Giarelis, N. Kanakaris and N. Karacapilidis: An innovative graph-based approach to advance feature selection from multiple textual documents. In: Maglogiannis I., Iliadis L., Pimenidis E. (eds), Artificial Intelligence Applications and Innovations - Proceedings of the 16th International Conference on Artificial Intelligence Applications and Innovations (AIAI 2020), Halkidiki, Greece, June 5-7, 2020, Springer, Cham, IFIP Advances in Information and Communication Technology, Vol 583, pp. 96-106, doi: 10.1007/978-3-030-49161-1_9. ([link](#))
15. N. Giarelis, N. Kanakaris and N. Karacapilidis: On a novel representation of multiple textual documents in a single graph. In: I. Czarnowski, R.J. Howlett and L.C.Jain (eds.), Intelligent Decision Technologies - Proceedings of the 12th KES International Conference on Intelligent Decision Technologies (KES-IDT 2020), Split, Croatia, June 17-19, 2020,

Springer, Singapore, Smart Innovation, Systems and Technologies, Vol. 193, pp. 105-115, doi: 10.1007/978-981-15-5925-9_9. ([link](#))

16. N. Kanakaris, N. Karacapilidis and G. Kournetas: On the exploitation of textual descriptions for a better-informed task assignment process. In: Proceedings of the 9th International Conference on Operations Research and Enterprise Systems (ICORES 2020), Valletta, Malta, February 22-24, 2020, Science and Technology Publications, pp. 304-310, DOI: 10.5220/0009151603040310. ([link](#))

17. N. Kanakaris, N. Karacapilidis, G. Kournetas and A. Lazanas: Combining Machine Learning and Operations Research Methods to Advance the Project Management Practice. In: Parlier G., Liberatore F., Demange M. (eds), Operations Research and Enterprise Systems. Communications in Computer and Information Science, Vol. 1162. Springer, Cham, 2020, pp. 135-155, DOI: 10.1007/978-3-030-37584-3_7. ([link](#))

18. A. Kanterakis, G. Iatraki, K. Pityanou, L. Koumakis, N. Kanakaris, N. Karacapilidis and G. Potamias: Towards Reproducible Bioinformatics: The OpenBio-C Scientific Workflow Environment. In: Proceedings of the 19th IEEE International Conference on Bioinformatics and Bioengineering (BIBE 2019), Athens, Greece, October 28-30, 2019, pp. 221-226. ([link](#))

19. N. Kanakaris, N. Karacapilidis and A. Lazanas: On the advancement of Project Management through a flexible integration of Machine Learning and Operations Research tools. In: Proceedings of the 8th International Conference on Operations Research and Enterprise Systems (ICORES 2019), Prague, Czech Republic, February 19-21, 2019, Science and Technology Publications, pp. 362-369, DOI: 10.5220/0007387103620369. ([link](#))